

Lambdachine: A Virtual Machine for Haskell

Thomas Schilling
Erudify

FP Afternoon 2013, Zürich, Switzerland

Introduction

- Haskell is a nice language and some very nice tools (GHCi, ThreadScope, criterion), but could be better.
- Profiling requires recompiling your program and all the libraries it depends on!
- Libraries are distributed in source form and compilation can be quite slow.
- Solution: virtual machine
 - Bonus: Send code over network?
 - Bonus: JIT compiler may even provide better performance


Option 1: Use JVM / CLR

- Pro: 100s to 1000s of person years of implementation effort
- Pro: Battle-tested and available on many platforms (at least in some form)
- Pro: Interface with huge number of existing libraries
- Con: Designed with different usage pattern in mind -- Haskell is very different in a number of aspects. GHC's runtime system takes advantage of Haskell's idiosyncrasies in a number of ways.
- Con: Interfacing with existing language requires mapping of types which is often awkward.


Option 2: Build custom VM

- Lots of effort, though it may be possible to reuse some of GHC's runtime system.
- Features can be tuned for executing Haskell.
- May also be a good platform for other (statically typed) functional languages.
- Use ideas and some code from open source projects: PyPy, V8, Mono, LuaJIT



Trace-based Just-in-time compilation



Trace-based Just-in-time compilation



Trace-based Just-in-time compilation




scope of optimisation

Trace-based Just-in-time Compilation


- Tracing JITs have recently found their way into the programming language mainstream: Tracemonkey (Javascript), LuaJIT 2, Android's Dalvik VM, PyPy (Python), others: SPUR (CIL/.net)
- Most of these languages are dynamically typed.
- Great for dynamic languages - very large static control flow graphs (due to runtime type checks).
- A trace-based JIT creates a specialised monomorphic version for each frequently executed path.
- Simple and quick compiler, thus short warm-up time.

Haskell is dynamic at runtime, too!

Thunks



Overloading



Thunks and Specialisation

- `map :: (a -> b) -> [a] -> [b]`
 - `map (+1) :: [Int] -> [Int]`
 - `map (+1) (generateList ...)`
 - `sum (map (+1) (upto 1 100))`
-
- Supercompilation is not yet ready for production (progress seems stalled?)

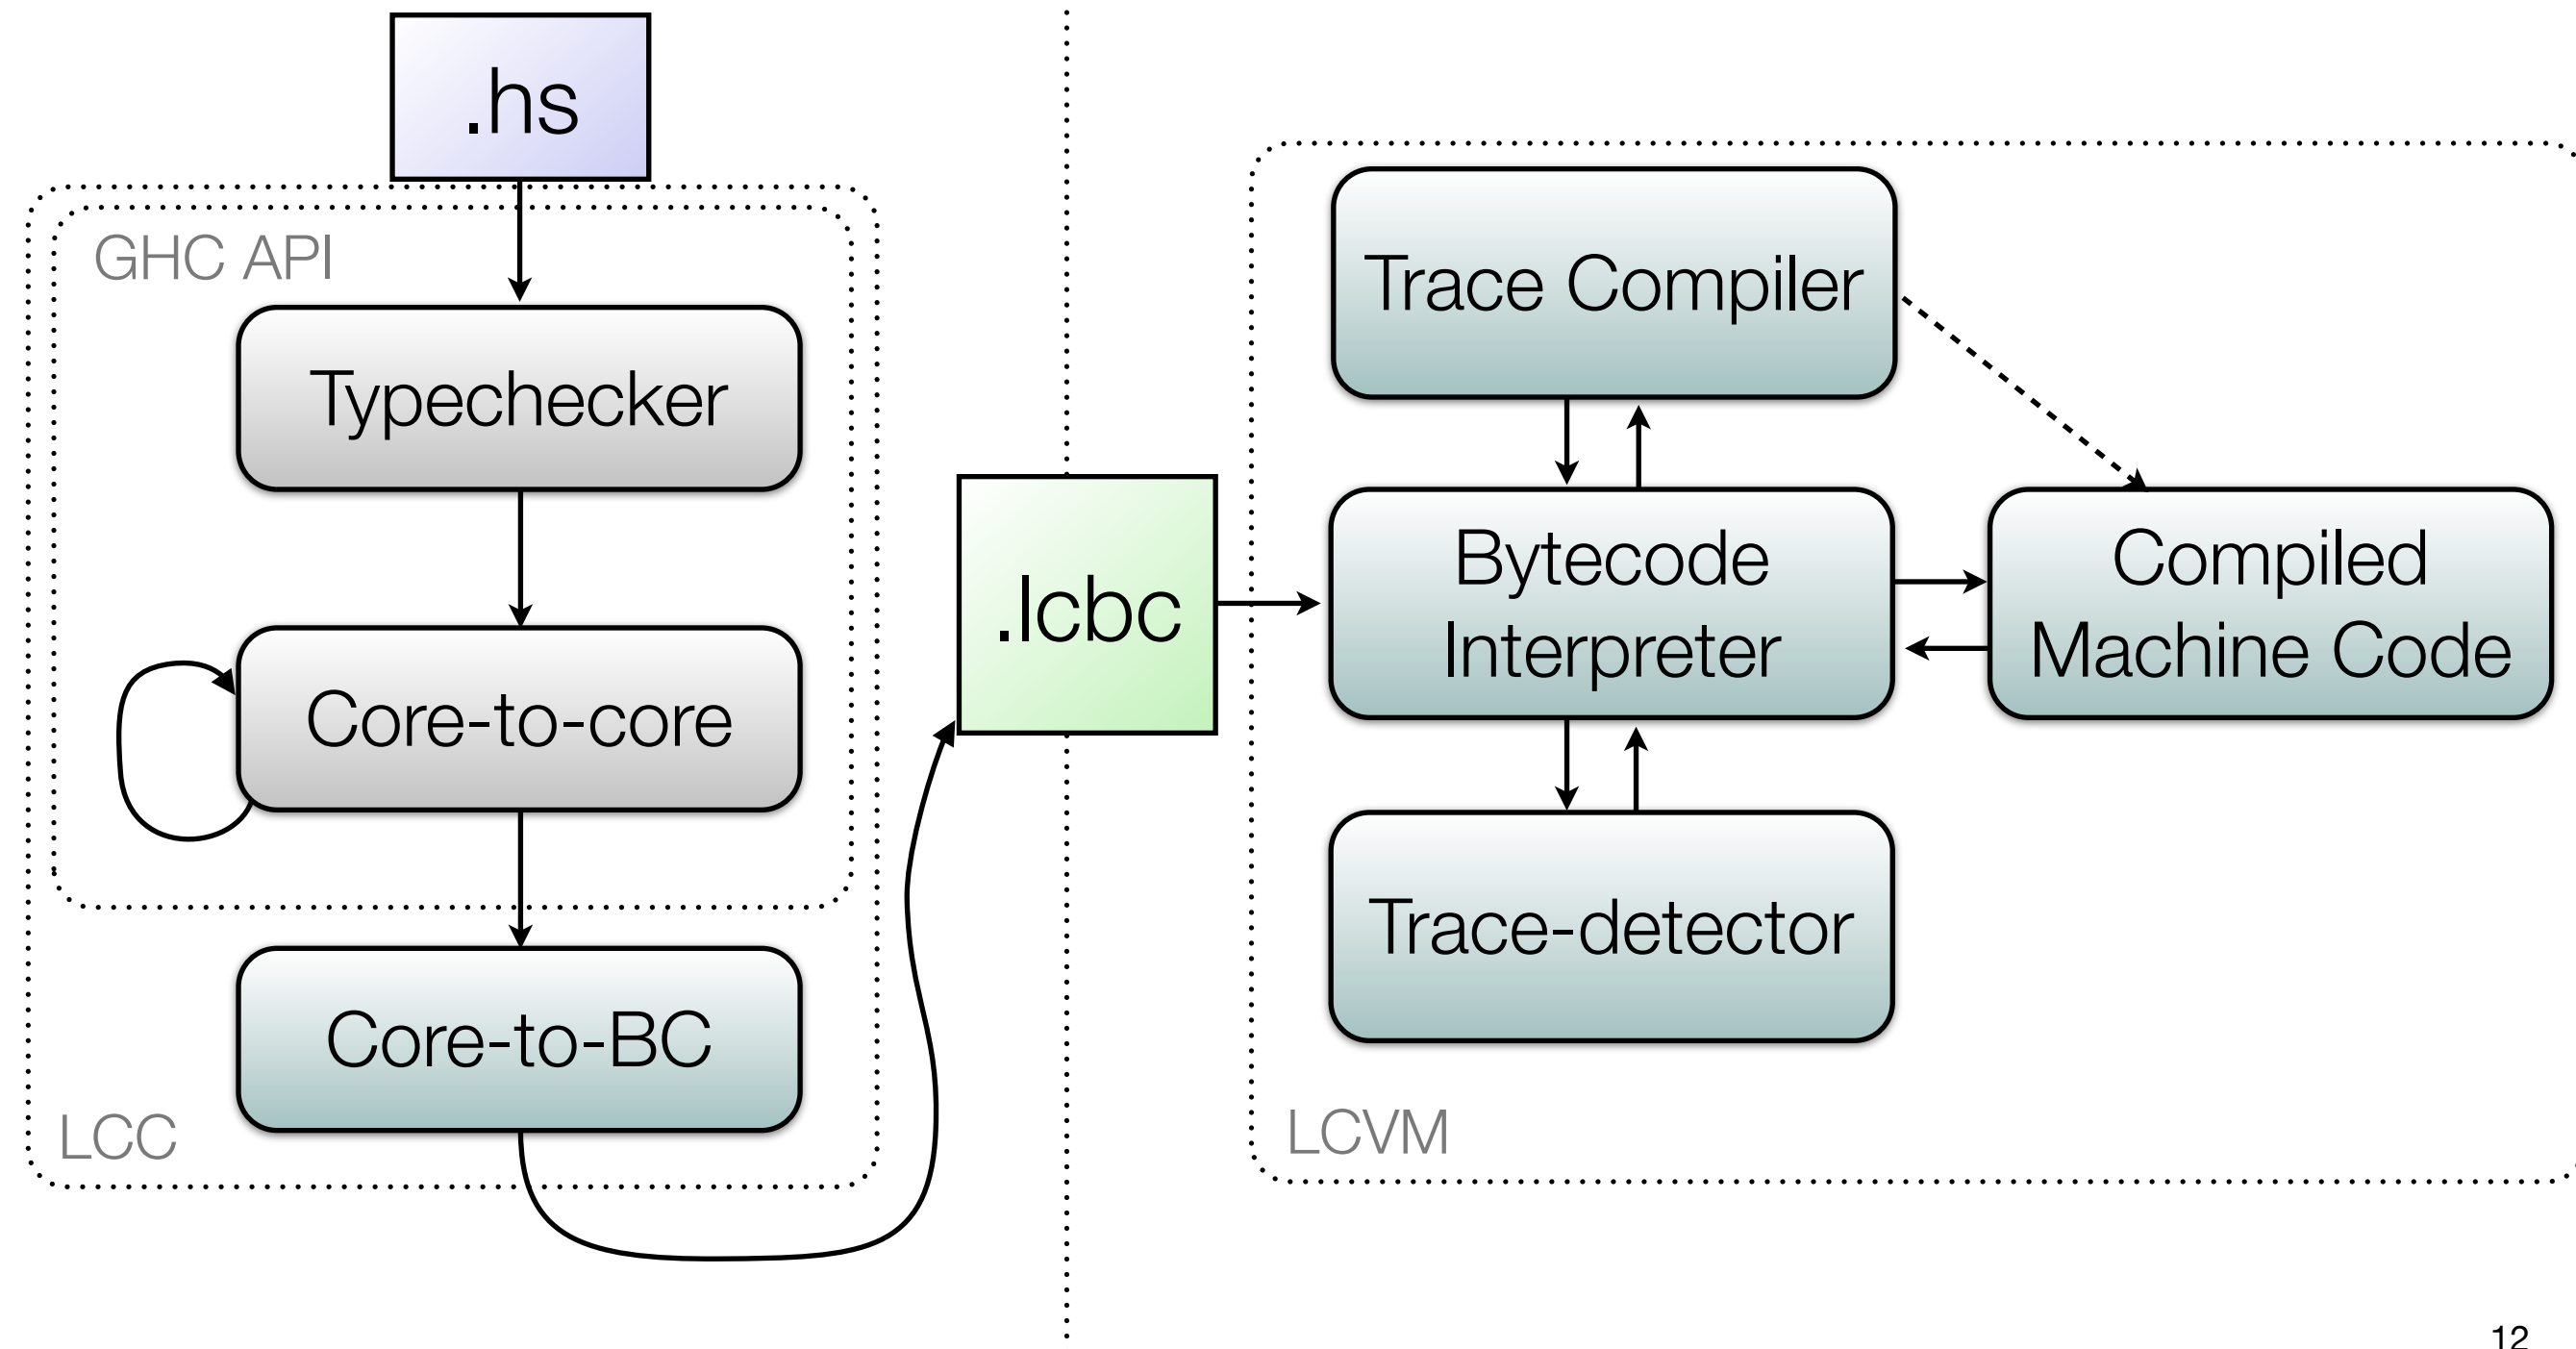
Overloading

- `add :: Num a => a -> a -> a`
- `add :: DNum a -> a -> a -> a`
- `add d x y = (plus d) x y`
- `addInt :: Int -> Int -> Int`
`addInt (I x) (I y) = I (primIntAdd x y)`

Lambdachine

Compilation Time

Execution Time



An Example

- `upto :: Int -> Int -> [Int]`
`upto lo hi =`
 `if lo > hi then [] else lo : upto (lo + 1) hi`
- `sum :: Int -> [Int] -> Int`
`sum !acc l = case l of`
 `[] -> acc`
 `(x:xs) -> sum (acc + x) xs`
- `test = print (sum 0 (upto 1 100))`

- `upto :: Int -> Int -> [Int]`
`upto lo hi =`
`if lo > hi then [] else lo : upto (lo + 1) hi`
- `sum :: Int -> [Int] -> Int`
`sum !acc l = case l of`
`[] -> acc`
`(x:xs) -> sum (acc + x) xs`
- `sum 55 (upto 11 100)`
`case (upto 11 100) of ...`
`case (upto 11 100) of ...`
`case (if 11 > 100 then ...) of ...`
`case (if False then ...) of ...`
`case (11 : upto 12 100) of ...`
`case (11 : upto 12 100) of ...`
`sum (55 + 11) (upto 12 100)`
`sum 66 (upto 12 100)`

start:

```
Obj *acc = base[0]; // load arg0
guard (info(acc) == Int);
Obj *l = base[1]; // load arg1
guard (info(l) == upto_thunk)
    // Enter thunk
int lo = l[1]; // load free var0
int hi = l[2]; // load free var1
guard (lo <= hi);
Obj *y = new I#(lo);
int lo2 = lo + 1;
Obj *ys = new upto_thunk(lo2, hi);
Obj *res = new Cons(y, ys);
update(l, res);
    // Return to sum
guard (info(res) == Cons);
Obj *x = res[0];
Obj *xs = res[1];
int x_u = x[0];
int acc_u = acc[0];
int acc_u2 = acc_u + x_u;
Obj *acc2 = new I#(acc_u2);
base[0] = acc2
base[1] = xs
goto start;
```

- upto :: Int# -> Int# -> [Int]
upto lo hi =
 if lo ># hi then [] else
 I# lo : upto (lo +# 1#) hi
- sum :: Int -> [Int] -> Int
sum !acc l = case l of
 [] -> acc
 (x:xs) -> sum (acc + x) xs

start:

```
Obj *acc = base[0]; // load arg0
guard (info(acc) == Int);
Obj *l = base[1]; // load arg1
guard (info(l) == upto_thunk)
    // Enter thunk
int lo = l[1]; // load free var0
int hi = l[2]; // load free var1
guard (lo <= hi);
Obj *y = new I#(lo);
```

- upto :: Int# -> Int# -> [Int]
upto lo hi =
 if lo ># hi then [] else
 I# lo : upto (lo +# 1#) hi
- sum :: Int -> [Int] -> Int
sum !acc l = case l of

A few microseconds later ...

xs

```
// Return to sum
guard (info(res) == Cons);
Obj *x = res[0];
Obj *xs = res[1];
int x_u = x[0];
int acc_u = acc[0];
int acc_u2 = acc_u + x_u;
Obj *acc2 = new I#(acc_u2);
base[0] = acc2
base[1] = xs
goto start;
```




```

start:
  Obj *acc = base[0];
  guard (info(acc) == Int);
  Obj *l = base[1];
  guard (info(l) == upto_thunk)
  int lo = l[1];
  int hi = l[2];
  guard (lo <= hi);
  Obj *y = new I#(lo);
  int lo2 = lo + 1;
  Obj *ys = new upto_thunk(lo2, hi);
  Obj *res = new Cons(y, ys);
  int acc_u = acc[0];
  int acc_u2 = acc_u + lo;
  Obj *acc2 = new I#(acc_u2);
loop:
  guard (lo2 <= hi);
  lo2 = lo2 + 1;
  acc_u2 = acc_u2 + lo2;
  goto loop;

```


Thunks and Updates

sum 55 (upto 11 100)




Thunks and Updates

sum 55 (11 : upto 12 100)




Thunks and Updates

sum 55 (11 : upto 12 100)




Thunks and Updates

sum 66 (upto 12 100)




Thunks and Updates

sum 66 (12 : upto 13 100)



Thunks and Updates

sum 78 (upto 13 100)



Benchmarks (micro)

Benchmark	GHC -O2	GHC -O2 + JIT
SumFromTo1	2.31s	1.80s
SumFromTo2	4.02s	2.32s
SumSquare	2.45s	2.35s
SumStream	0.24s	1.11s
Tak	1.01s	0.84s

Benchmarks (small)

Benchmark	GHC -O2	GHC -O2 + JIT
WheelSieve2	0.34s	0.59s
Boyer	0.75s	0.74s
Constraints	0.88s	0.89s
Circsim	1.71s	3.49s
Lambda	0.74s	1.06s

Conclusions & Ongoing Work

- Trace selection is tricky. Haskell's control flow graphs are large messy. Preferring tail-recursive loops could help.
- How can we detect (cheaply) when an update can be omitted? -- Some recent work by SPJ et al
- Only a subset of the Prelude currently supported.
- Single-threaded and very simple garbage collector -- need to integrate GHC runtime system (scheduler, generational GCs, sparks, STM, ...)

Questions?
