


Haxl: Haskell at Facebook


Simon Marlow
Jon Coens
Louis Brandy
Bartosz Nitka
Jon Purdy
Aaron Roth
& others





What's in this talk

- The Haxl project: ~12 months later, where are we
- Haxl published at ICFP'14!
- Haxl open source release!
 - walking through an example data source










No!





No!

How does Sigma know what's spam?

- FXL

```
JaffaCakeSpam =  
  MessageContains("Jaffa Cakes") &&  
  Let  
    LikesJaffaCakes(X) = Likes(JaffaCakes,X)  
  In  
    Length(Filter(LikesJaffaCakes, FriendsOf(SourceId))) < 3
```

- + machine learning.

```
SpamMessage = SpamScore > 0.99
```

- Where do the inputs to the ML come from?
 - FXL expressions.

What can you do in FXL?

- Fetch data from the Facebook graph:

```
FriendCount(uid) = AssocCountByType(uid, AssocFriends)
```

- Fetch data from any of the other 18 data sources
- Run machine learning classifiers
- Perform simple computations

```
RatioFriendsSourceIdOver20 =  
  If FriendCountSourceId > 0  
    Then Ratio(CountFriendsSourceIdOver(20),  
               FriendCountSourceId)  
  Else 0.0;
```


What's good about FXL?

- Clean syntax
 - SI engineers concentrate on fighting spam, not the language
- Static typing
 - cannot push type-incorrect code
- We can push changes *fast*
 - a couple of minutes from commit to production

What's not so good?

What's not so good?

- Limited abstractions

We're building larger systems in FXL now

What's not so good?

- Limited abstractions

We're building larger systems in FXL now

- Design quirks and hysterical raisins

What's not so good?

- Limited abstractions

We're building larger systems in FXL now

- Design quirks and hysterical raisins

Static typing is limited

- Only a few types: int, double, string, vector, map, JSON
- No user-defined types
- Type system doesn't catch as many errors as it could

What's not so good?

- Limited abstractions

We're building larger systems in FXL now

- Design quirks and hysterical raisins

Static typing is limited

- Only a few types: int, double, string, vector, map, JSON
 - No user-defined types
 - Type system doesn't catch as many errors as it could
- Slow (it's an interpreter)

Why are we switching to Haskell?

- Expressivity
- Learning resources available
- Lots of libraries
- Faster (it's compiled)
- Better implementation (error messages etc.)
- Chance to redesign the whole system
 - guaranteed replayability

Technical challenges

1. Implicit concurrency
 2. Implement all the FXL functionality in Haskell
 3. Translate all the FXL code
 4. Figure out how to compile+push all the Haskell code to all of the machines in a few minutes
- Status summary: we're mostly done with 1,2,3 and experimenting with a solution for 4.
 - Now: testing, bug fixing and optimisation.

Implicit concurrency

- In FXL you can write this:

```
NumCommonFriends(x, y) =  
    Length(Intersect(FriendsOf(x), FriendsOf(y)));
```

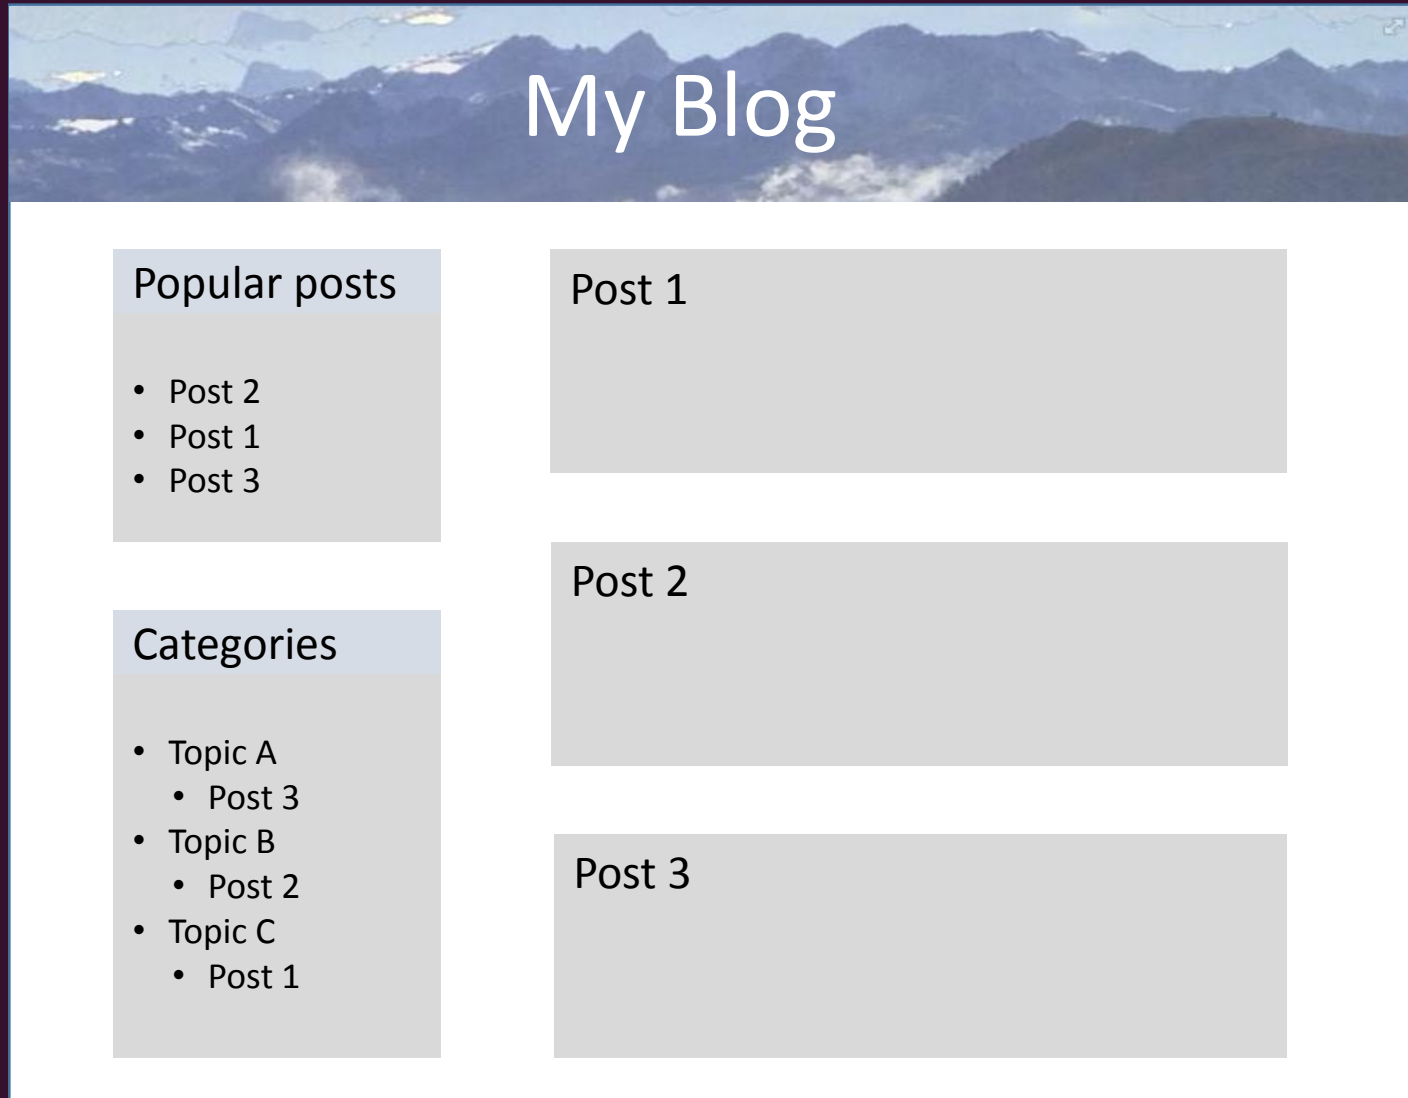
- And Sigma automatically batches the two requests together.
- With existing languages & frameworks you have to specify the concurrency explicitly...

- e.g using Haskell asyncs:


```
do
    ax <- async (friendsOf x)
    ay <- async (friendsOf y)
    fx <- wait ax
    fy <- wait ay
    return (length (intersect fx fy))
```

- Too verbose
- Prone to false dependencies

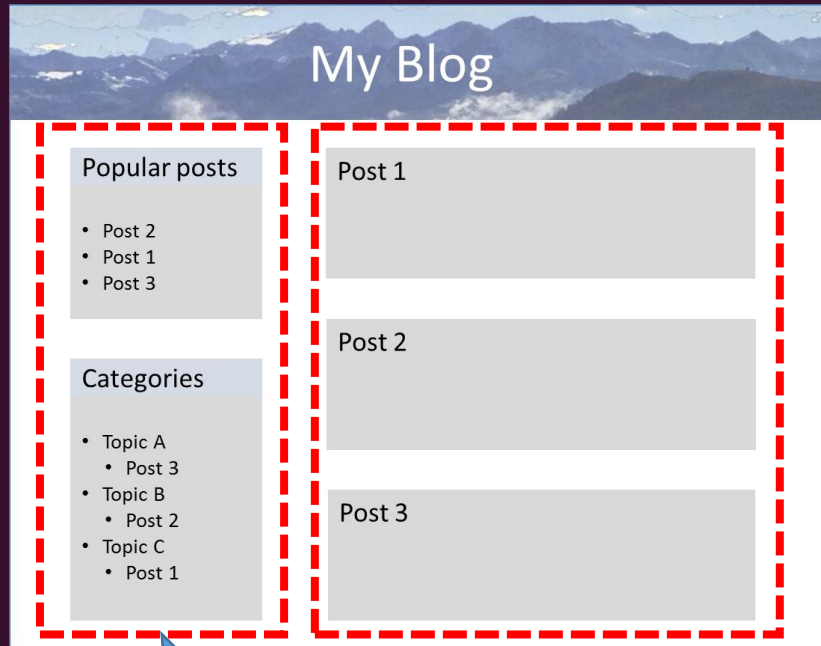
Larger example: a blog server



Another example: a blog server

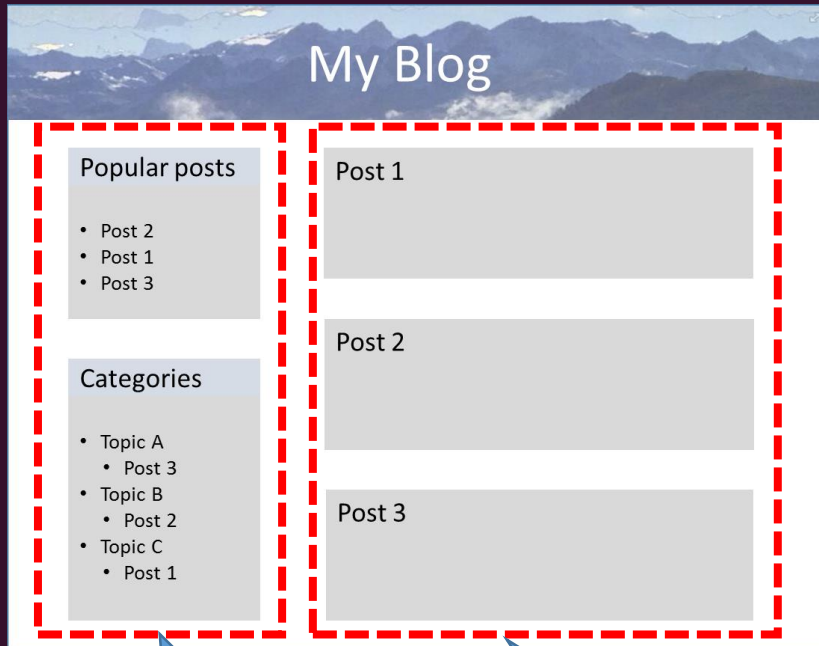


Another example: a blog server



We want modular code –
code each pane
independently

Another example: a blog server



We want modular code –
code each pane
independently

We want it to execute
efficiently.

- Concurrent data fetches
- No repeated data fetches

```
data PostId      -- identifies a post
data Date        -- a calendar date
data PostContent -- the content of a post

data PostInfo = PostInfo
  { postId      :: PostId
  , postDate    :: Date
  , postTopic   :: String
  }

-- data-fetching operations
getPostIds      :: Haxl [PostId]
getPostInfo     :: PostId -> Haxl PostInfo
getPostContent  :: PostId -> Haxl PostContent
getPostViews    :: PostId -> Haxl Int

-- rendering functions
renderPosts     :: [(PostInfo,PostContent)] -> Html
renderPage      :: Html -> Html -> Html
...
```

```
blog :: Haxl Html
blog = renderPage <$> leftPane <*> mainPane

getAllPostsInfo :: Haxl [PostInfo]
getAllPostsInfo = mapM getPostInfo =<< getPostIds

mainPane :: Haxl Html
mainPane = do
  posts <- getAllPostsInfo
  let ordered =
    take 5 $
    sortBy (flip (comparing postDate)) posts
  content <- mapM (getPostContent . postId) ordered
  return $ renderPosts (zip ordered content)
```



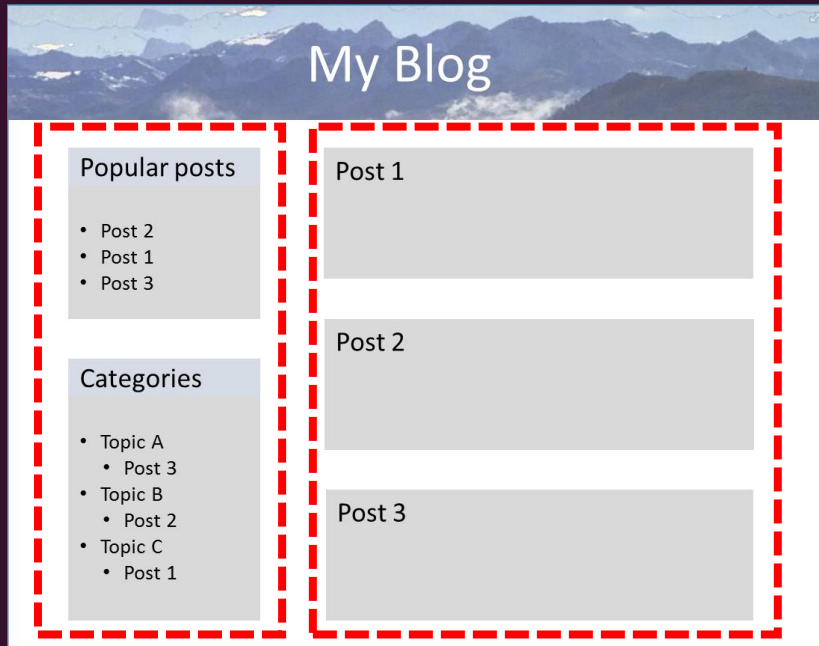
```
leftPane :: Haxl Html
leftPane = renderSidePane <$> popularPosts <*> topics

getPostDetails :: PostId -> Haxl (PostInfo, PostContent)
getPostDetails pid =
  (,) <$> getPostInfo pid <*> getPostContent pid


popularPosts :: Haxl Html
popularPosts = do
  pids <- getPostIds
  views <- mapM getPostViews pids
  let ordered =
    take 5 $ map fst $
    sortBy (flip (comparing snd)) (zip pids views)
  content <- mapM getPostDetails ordered
  return $ renderPostList content
```

```
topics :: Haxl Html
topics = do
  posts <- getAllPostsInfo
  let topiccounts =
    Map.fromListWith (+) [ (postTopic p, 1) | p <- posts ]
  return $ renderTopics topiccounts
```

So how did we do?

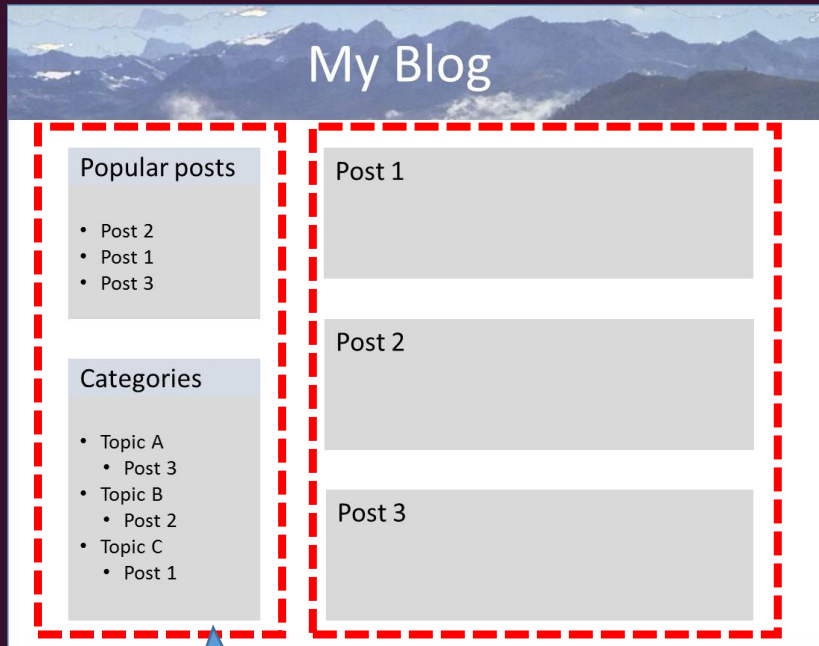


So how did we do?



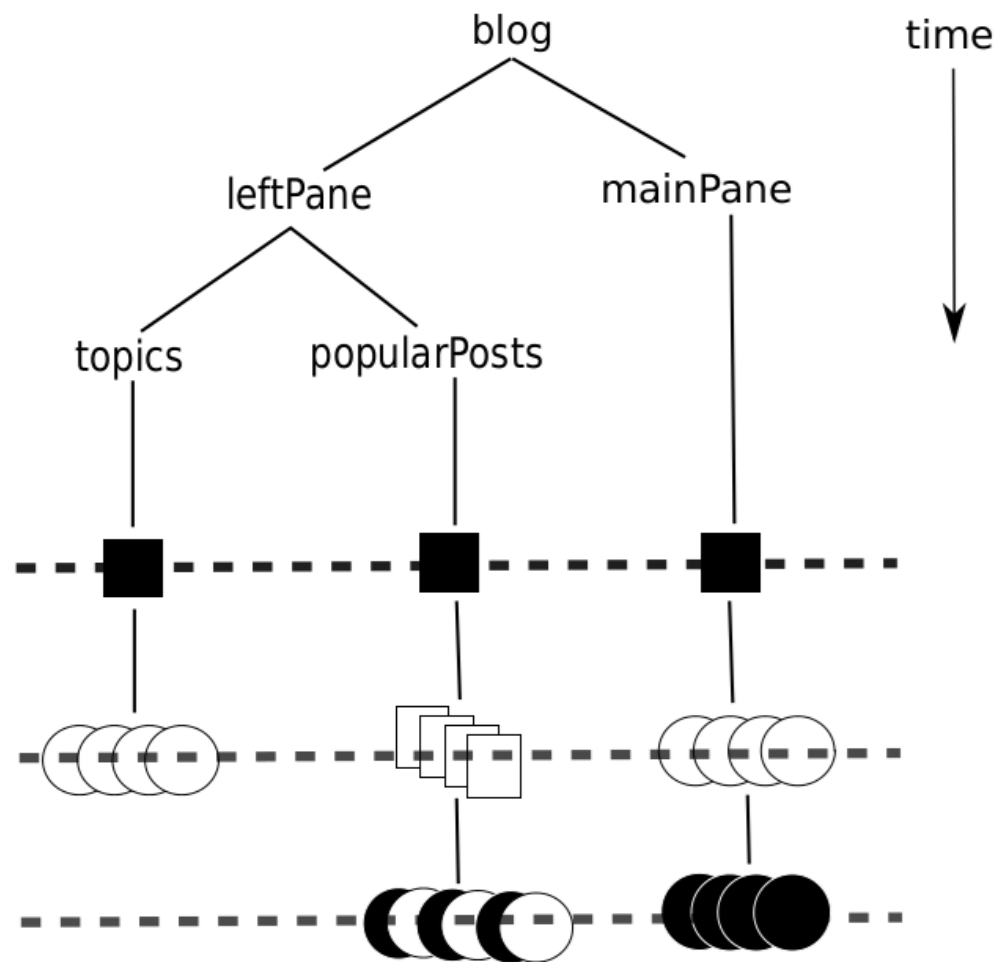
Code is clean and modular

So how did we do?



Code is clean and modular

But does it execute efficiently?



Demo

Implementation

- Start with a concurrency monad

```
data Result a

    -- we're finished, here's the result
= Done a

    -- the computation blocked...
| Blocked
    (Seq BlockedRequest) -- requests to perform
    (Haxl a)              -- continuation
```

Start with a concurrency monad

```
data Result a
  = Done a
  | Blocked (Seq BlockedRequest) (Haxl a)

newtype Haxl a = Haxl { unHaxl :: IO (Result a) }

instance Monad Haxl where
  return a = Haxl $ return (Done a)

Haxl m >>= k = Haxl $ do
  r <- m
  case r of
    Done a      -> unHaxl (k a)
    Blocked br c -> return (Blocked br (c >>= k))
```

Add an Applicative instance

```
instance Applicative Haxl where
  pure = return

Haxl f <*> Haxl x = Haxl $ do
  f' <- f
  x' <- x
  case (f',x') of
    (Done g,      Done y      ) -> return (Done (g y))
    (Done g,      Blocked br c ) -> return (Blocked br (g <$> c))
    (Blocked br c, Done y      ) -> return (Blocked br (c <*> return y))
    (Blocked br1 c, Blocked br2 d) -> return (Blocked (br1 <> br2) (c <*> d))
```

Fetching Data

```
dataFetch :: Request a -> Haxl a
```

GADT

Type parameter is
the result type of
the request

Fetching Data

```
dataFetch :: Request a -> Haxl a
```

GADT

Type parameter is
the result type of
the request

```
data Request a where
  FetchPosts      :: Request [PostId]
  FetchPostInfo   :: PostId -> Request PostInfo
  FetchPostContent :: PostId -> Request PostContent
  FetchPostViews  :: PostId -> Request Int
```

Fetching Data

```
dataFetch :: Request a -> Haxl a
```

GADT

Type parameter is
the result type of
the request

```
data Request a where
  FetchPosts      :: Request [PostId]
  FetchPostInfo   :: PostId -> Request PostInfo
  FetchPostContent :: PostId -> Request PostContent
  FetchPostViews  :: PostId -> Request Int
```

```
data FetchStatus a = NotFetched | FetchSuccess a
```

```
data BlockedRequest =
  forall a . BlockedRequest (Request a) (IORef (FetchStatus a))
```

- We can implement dataFetch:

```
dataFetch :: Request a -> Haxl a
dataFetch request = Haxl $ do
  box <- newIORef NotFetched
  let br = BlockedRequest request box
  let cont = Haxl $ do
    FetchSuccess a <- readIORef box
    return (Done a)
  return (Blocked (singleton br) cont)
```

- To fetch data, we need

```
fetch :: [BlockedRequest] -> IO ()
```

Application-specific data-fetching function.

Batches multiple requests, uses concurrency, etc.

- To run a computation to completion, we need a loop:

```
runHaxl :: Haxl a -> IO a
runHaxl (Haxl h) = do
  r <- h
  case r of
    Done a -> return a
    Blocked br cont -> do
      fetch (toList br)
      runFetch cont
```

- Done!

We also want caching

- Reader monad passes an **IORef DataCache** around
- Complication:
 - cache maps **Request a** to **a**
 - can't do this with **Data.Map** alone

```
newtype DataCache =  
  DataCache (forall a . Map (Request a) (IORef (FetchStatus a)))
```

There is no Fork: an Abstraction for Efficient, Concurrent, and Concise Data Access

Simon Marlow

Facebook
smarlow@fb.com

Louis Brandy

Facebook
ldbrandy@fb.com

Jonathan Coens

Facebook
jon.coens@fb.com

Jon Purdy

Facebook
jonp@fb.com

Abstract

We describe a new programming idiom for concurrency, based on Applicative Functors, where concurrency is implicit in the Applicative `<*>` operator. The result is that concurrent programs can be written in a natural applicative style, and they retain a high degree of clarity and modularity while executing with maximal concurrency. This idiom is particularly useful for programming against external data sources, where the application code is written without the use of explicit concurrency constructs, while the implementation is able to batch together multiple requests for data from the same source, and fetch data from multiple sources concurrently. Our abstraction uses a cache to ensure that multiple requests for the same data return the same result, which frees the programmer from having to arrange to fetch data only once, which in turn leads to greater modularity.

While it is generally applicable, our technique was designed with a particular application in mind: an internal service at Facebook that identifies particular types of content and takes actions based on it. Our application has a large body of business logic that fetches data from as many as 15 different external sources. The framework described in this paper enables the business logic to execute efficiently by automatically fetching data concurrently; we present some preliminary results.

efficiency in this setting: accessing multiple remote data sources efficiently requires *concurrency*, and that normally requires the programmer to intervene and program the concurrency explicitly.


When the business logic is only concerned with *reading* data from external sources and not *writing*, the programmer doesn't care about the order in which data accesses happen, since there are no side-effects that could make the result different when the order changes. So in this case the programmer would be entirely happy with not having to specify either ordering or concurrency, and letting the system perform data access in the most efficient way possible. In this paper we present an embedded domain-specific language (EDSL), written in Haskell, that facilitates this style of programming, while automatically extracting and exploiting any concurrency inherent in the program.

Our contributions can be summarised as follows:

- We present an **Applicative** abstraction that allows implicit concurrency to be extracted from computations written with a combination of **Monad** and **Applicative**. This is an extension of the idea of concurrency monads [9], using Applicative `<*>` as a way to introduce concurrency (Section 4). We then develop the idea into an abstraction that supports concurrent access to remote data (Section 5), and failure (Section 8).
- We show how to add a *cache* to the framework (Section 6).

But...

- The Request type was wired into the monad
- How can we make the monad independent of the data source(s)?



- Core code includes the monad, caching support etc.
- Core is *generic*: no data sources built-in

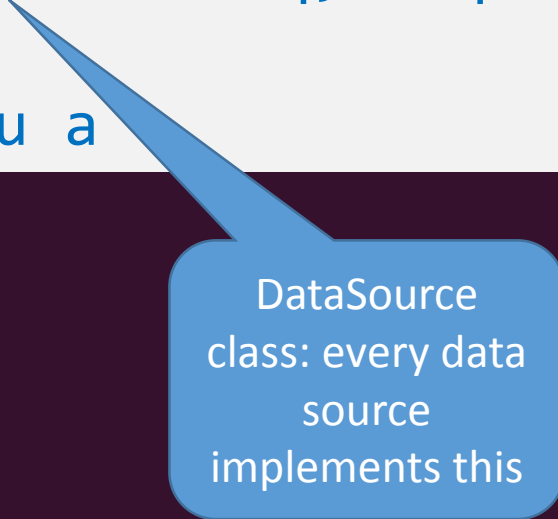
```
dataFetch :: Request a -> Haxl a
```

~~dataFetch :: Request a -> Haxl a~~

```
dataFetch :: (DataSource u req, Request req a)
           => req a
           -> GenHaxl u a
```

```
dataFetch :: (DataSource u req, Request req a)
           => req a
           -> GenHaxl u a
```

```
dataFetch :: (DataSource u req, Request req a)
           => req a
           -> GenHax1 u a
```



DataSource
class: every data
source
implements this


```
dataFetch :: (DataSource u req, Request req a)
           => req a
           -> GenHax1 u a
```

DataSource
class: every data
source
implements this

Request class:
just Eq,
Hashable,
Typeable, Show

```
dataFetch :: (DataSource u req, Request req a)
           => req a
           -> GenHax1 u a
```

DataSource
class: every data
source
implements this

Request class:
just Eq,
Hashable,
Typeable, Show

User state –
passed around,
can be accessed
by data sources

DataSource walk-through

- We'll walk through constructing a complete data source
- We'll make a data source for the Facebook Graph API
 - web API for querying the Facebook Graph
 - using Felipe Lessa's **fb** package to do the real work
 - our data source will perform requests concurrently up to a maximum number of threads

- Start with the request type:

```
data FacebookReq a where
  GetObject      :: Id -> FacebookReq Object
  GetUser        :: UserId -> FacebookReq User
  GetUserFriends :: UserId -> FacebookReq [Friend]
deriving Typeable
```



GADT, as before

- We also need some boilerplate:

```
deriving instance Eq (FacebookReq a)
deriving instance Show (FacebookReq a)

instance Show1 FacebookReq where show1 = show

instance Hashable (FacebookReq a) where ...
```

- A data source has some state:

```
instance StateKey FacebookReq where
  data State FacebookReq =
    FacebookState
      { credentials :: Credentials
      , userAccessToken :: UserAccessToken
      , manager :: Manager
      , numThreads :: Int
      }
```

API keys

HTTP connection
manager

Concurrency
control

- Initialise the state:

```
initGlobalState
  :: Int
  -> Credentials
  -> UserAccessToken
  -> IO (State FacebookReq)

initGlobalState threads creds token = do
  manager <- newManager tlsManagerSettings
  return FacebookState
    { credentials = creds
    , manager = manager
    , userAccessToken = token
    , numThreads = threads
    }
```

- nothing surprising there.

- Make an instance of DataSource

```
class DataSourceName req where
  dataSourceName :: req a -> Text

class (DataSourceName req, StateKey req, Show1 req)
  => DataSource u req where
  fetch
    :: State req
    -> Flags
    -> u
    -> [BlockedFetch req]
    -> PerformFetch
```

```
instance DataSourceName FacebookReq where
  dataSourceName _ = "Facebook"

instance DataSource u FacebookReq where
  fetch = facebookFetch
```

- Implement fetch

```
data PerformFetch
  = SyncFetch  (IO ())
  | AsyncFetch (IO () -> IO ())
```

```
facebookFetch
  :: State FacebookReq
  -> Flags
  -> ()
  -> [BlockedFetch FacebookReq]
  -> PerformFetch
```

IO to do while
the requests
are in progress

```
facebookFetch FacebookState{..} _flags _user bfs =
  AsyncFetch $ \inner -> do
    sem <- newQSem numThreads
    asyncs <- mapM (async . fetchAsync credentials manager
                                     userAccessToken sem) bfs

    inner
    mapM_ wait asyncs
```


- Implement fetchAsync

```
fetchAsync
```

```
  :: Credentials -> Manager -> UserAccessToken -> QSem  
  -> BlockedFetch FacebookReq  
  -> IO ()
```

```
fetchAsync creds manager tok sem (BlockedFetch req rvar) =  
  bracket_ (waitQSem sem) (signalQSem sem) $ do
```

```
    e <- Control.Exception.try $  
        runResourceT $  
        runFacebookT creds manager $  
        fetchReq tok req
```

```
    case e of  
      Left ex -> putFailure rvar (ex :: SomeException)  
      Right a  -> putSuccess rvar a
```

- fetchReq maps FacebookReq to FacebookT computations

```
fetchReq
  :: UserAccessToken
  -> FacebookReq a
  -> FacebookT Auth (ResourceT IO) a

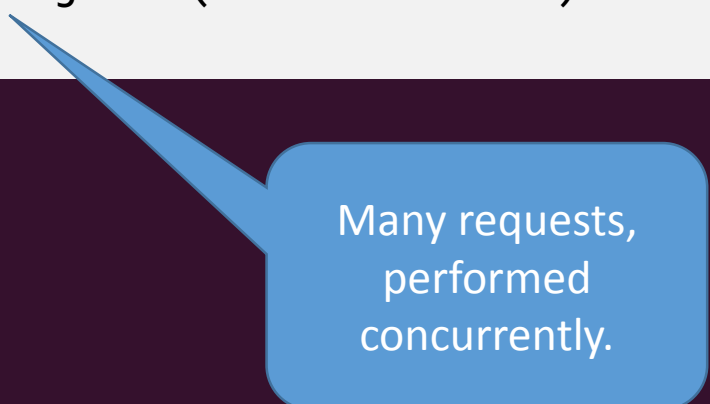
fetchReq tok (GetObject (Id id)) =
  getObject ("/" <> id) [] (Just tok)

fetchReq _tok (GetUser id) =
  getUser id [] Nothing

fetchReq tok (GetUserFriends id) = do
  f <- getUserFriends id [] tok
  source <- fetchAllNextPages f
  source $$ consume
```

- Example

```
main :: IO ()
main = do
  (creds, access_token) <- getCredentials
  facebookState <- initGlobalState 10 creds access_token
  env <- initEnv (stateSet facebookState stateEmpty) ()
  r <- runHaxl env $ do
    likes <- getObject "me/likes"
    mapM getObject (likeIds likes)
  print r
```



Many requests,
performed
concurrently.

Back to our Haxl project...

- But do people have to learn `<$>`, `<*>`, etc?

```
numCommonFriends x y =  
  length <$> (intersect <$> friendsOf x <*> friendsOf y)
```

- No, because this

```
numCommonFriends x y = do  
  fx <- friendsOf x  
  fy <- friendsOf y  
  return (length (intersect fx fy))
```

- can be silently translated to the Applicative form in the compiler
 - (not implemented yet)

- Going further, we could write a pre-processor from this:

```
numCommonFriends :: Haxl Int
numCommonFriends =
    length (intersect (friendsOf sourceId) (friendsOf targetId))
```

- To this:

```
numCommonFriends :: Haxl Int
numCommonFriends
= length <$>
    (intersect <$>
        (join (friendsOf <$> sourceId)) <*>
        (join (friendsOf <$> targetId)))
```

- straightforward with **haskell-src-extends**

Dilemma: monads or no monads?

- Using a preprocessor
 - Advantages
 - Everything is monadic, but looks pure to the programmer
 - Easier to understand
 - Disadvantages
 - Can't write pure code
 - Hard to interpret error messages
 - Two languages adds complexity
- We decided not to go this route (for now)

Technical challenges

1. Implicit concurrency
2. Implement all the FXL functionality in Haskell
3. Translate all the FXL code
4. Figure out how to compile+push all the Haskell code to all of the machines in a few minutes

Implement all the FXL functionality in Haskell

- several data sources
- Each needs a Haskell/C++ FFI layer
- ~450 built-in functions
 - Ranging from easy (StrCmp) to really annoying (ParseActivityLog)
- We created lots of tasks
 - some done by the team
 - (we're working on our 4th iteration of the TAO layer)
 - others grabbed by interested people around Facebook: hack-a-month projects and bootcamp

- As of two weeks ago, we have everything implemented!

Technical challenges

1. Implicit concurrency
2. Implement all the FXL functionality in Haskell
3. Translate all the FXL code
4. Figure out how to compile+push all the Haskell code to all of the machines in a few minutes

Translate all the FXL code

- We have a *lot* of FXL code
 - impractical to translate it all by hand
- Wrote a translation tool
 - tricky bit is converting to do-syntax or Applicative where necessary, while keeping as much code as possible pure
- Auto-translated code will become the source
 - Try to produce readable code

Technical challenges

1. Implicit concurrency
2. Implement all the FXL functionality in Haskell
3. Translate all the FXL code
4. Figure out how to compile+push all the Haskell code to all of the machines in a few minutes

Compile time

- At first, compiling the whole translated codebase took ~30 mins
- (FXL push currently takes ~2 mins)




Reducing compile time

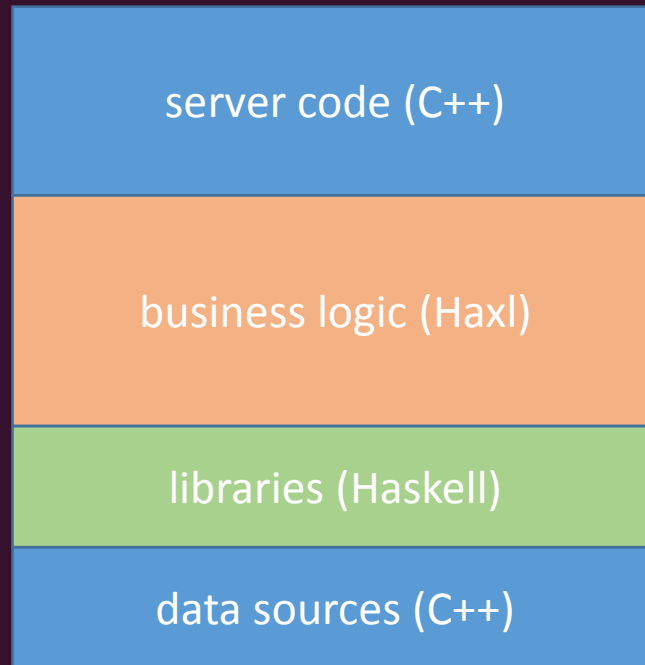
- Long laborious process to impose a sensible module structure
 - FXL source files now form a DAG
 - Compilation has some parallelism now
- We have full compile down to ~5 mins
 - incremental compile usually much faster (~2 mins)

How to push to Sigma machines?


- We're experimenting with GlusterFS for distribution.
 - seems good: all machines get a new object in <1 min



Hot code swapping



Hot code swapping



Hot code swapping

- Keep serving requests while we load new code
- Use GHC's built-in linker
 - Had to modify it to unload code (shipped in GHC 7.8)
 - GC detects when it is safe to release old code
 - We can have multiple copies of the code running while existing requests drain

Status

- Call graph complete
- Full FXL codebase translated
- Next goals:
 - achieve 99% correctness (100% hard due to random effects)
 - get performance up par with FXL
 - experiment with running production traffic
 - open source...

Open Source! (coming next week)

The screenshot shows a GitHub repository page for 'facebook / Haxl'. The repository is private and has 12 stars, 2 forks, and 0 releases. It was created by Simon Marlow. The repository description is 'A Haskell library for efficient, concurrent, concise data access.' The repository contains 1 commit, 1 branch (master), and 0 contributors. The repository was created by Simon Marlow on Tuesday at 16:10. The repository contains the following files and folders: Haxl, example, tests, .gitignore, LICENSE, PATENTS, Setup.hs, TARGETS, haxl.cabal, and readme.md. The repository was created by Simon Marlow on Tuesday at 16:10. The repository contains the following files and folders: Haxl, example, tests, .gitignore, LICENSE, PATENTS, Setup.hs, TARGETS, haxl.cabal, and readme.md. The repository was created by Simon Marlow on Tuesday at 16:10.

Initial open source import

Simon Marlow authored Tuesday at 16:10 latest commit 7120357ec9

Haxl	Initial open source import	a day ago
example	Initial open source import	a day ago
tests	Initial open source import	a day ago
.gitignore	Initial open source import	a day ago
LICENSE	Initial open source import	a day ago
PATENTS	Initial open source import	a day ago
Setup.hs	Initial open source import	a day ago
TARGETS	Initial open source import	a day ago
haxl.cabal	Initial open source import	a day ago
readme.md	Initial open source import	a day ago

readme.md

Haxl

Haxl is a Haskell library that simplifies access to remote data, such as databases or web-based services. Haxl can automatically

- batch multiple requests to the same data source,
- request data from multiple data sources concurrently.

Code

Issues 0

Pull Requests 0

Wiki

Pulse

Graphs

Network

SSH clone URL

git@github.com:face

You can clone with [HTTPS](#), [SSH](#), or [Subversion](#).

Clone in Desktop

Download ZIP

Haxl: Haskell at Facebook

Simon Marlow
Jon Coens
Louis Brandy
Bartosz Nitka
Jon Purdy
Aaron Roth
& others
<your name here>



Questions?